

PRODUCT REQUIREMENTS DOCUMENT (PRD)

Sentient Self — AI-Powered Therapeutic Growth Platform

Version: 1.0 (Locked)

Status: Build-Ready

Audience: AI Coding Agent / Engineering Team

Platform: Web Application (Responsive)

1. PRODUCT OVERVIEW

1.1 Product Vision

Sentient Self is a **conversation-first, AI-guided therapeutic self-development platform** designed primarily for **recovering addicts**, and secondarily for **individuals seeking deep personal growth**.

The application simulates the *supportive flow of therapy* by prioritizing:

- Emotional validation
- Context gathering
- Meaning-making
- Optional structured exercises introduced only when appropriate

The system is **not medical care**, but includes **robust crisis detection and escalation safeguards**.

1.2 Target Audience

Primary

- Individuals in addiction recovery
- Users experiencing cravings, shame, relapse risk, emotional dysregulation

Secondary

- Individuals pursuing self-awareness, purpose, trauma processing
- Corporate wellness users
- Treatment-center clients

2. PLATFORM & DEPLOYMENT

Aspect	Specification
Platform	Web-only
Mobile Support	Fully responsive
Offline Mode	Not supported
Dark Mode	Required
Accessibility	No explicit requirements
Hosting	Cloud-based

3. AUTHENTICATION & ACCOUNTS

3.1 Authentication

- Required to use app
- Methods:
 - Email / password
 - Google OAuth
 - Apple OAuth

3.2 Account Types

1. **Personal Account**
2. **Institution-Provisioned User Account**
3. **Institution Admin Account**

Institution users may later convert to personal paid accounts.

4. MONETIZATION

4.1 Plans

Free Tier

- Unlimited conversational AI
- Unlimited journal entries
- 5 AI-guided exercises per month
- ~~✗~~ No AI journal analysis
- ~~✗~~ AI model limited to **DeepSeek R1**

Premium Tier — \$14.99/month

- Unlimited exercises
- AI journal analysis
- Full analytics dashboard
- Gamification system
- Access to all AI models

4.2 Institutional Licensing

- Monthly **per-seat** billing
- Stripe subscriptions & invoicing
- Admin seat management
- Aggregate analytics only (no content access)

5. SAFETY, LEGAL & CRISIS HANDLING

5.1 Medical Disclaimer

- Displayed during onboarding (required acknowledgment)
- Re-shown during crisis events
- App explicitly states it is ****not medical care****

5.2 Crisis Detection

Triggers include:

- Suicidal ideation
- Self-harm intent
- Imminent danger
- Severe mental health crisis language

5.3 Crisis Response (US-Only)

- Immediate interruption of conversation
- Transition to ****Crisis Mode****
- Display:
 - 988 Suicide & Crisis Lifeline
 - 911 for emergencies
- No exercises, reflection, or meaning-making

5.4 Institutional Crisis Alerts

- ****Identifiable alert sent to institution****
- Only for institution-provisioned users
- Explicit user consent required during onboarding
- No journal or chat content shared

6. CORE UX PHILOSOPHY

- Conversation is always the starting point
- Exercises are ****optional invitations****
- Users never browse or select exercises manually
- The AI introduces structure only when contextually appropriate
- User autonomy is non-negotiable

7. ONBOARDING ASSESSMENT (MVP)

7.1 Purpose

Establish an initial ****User Baseline Profile**** to prime AI context without overwhelming the user.

****Duration:**** ~5–8 minutes

****All questions skippable except consent****

7.2 Onboarding Flow

1. **Welcome & Safety Framing**
2. **Intent & Desires**
3. **Current Challenges**
4. **Addiction Context (Conditional)**
5. **Emotional Baseline**
6. **Support Preferences**
7. **Consent & Transparency**
8. **AI Warm Hand-Off**

7.3 User Baseline Profile Object

```
```json
UserBaselineProfile = {
 primary_intents: [],
 current_challenges: [],
 addiction_context: {
 substances: [],
 recovery_stage: ""
 },
 emotional_baseline: [],
 support_preferences: {
 tone: "",
 exercise_openness: ""
 },
 onboarding_completed_at: timestamp
}```
```

#### \*\*Rules\*\*

- Stored privately
- Used for AI context only
- Never treated as fixed truth
- Not shown verbatim to user

---

## ## 8. AI SYSTEM OVERVIEW

### ### 8.1 AI Models

#### \*\*Free Tier\*\*

- DeepSeek R1 only

#### \*\*Premium\*\*

- GPT-5.2
- Claude 4.5 Opus
- Gemini 3 Pro
- Grok 4.1

Users may:

- Set a global default
- Assign models per feature
- **✗** Cannot switch mid-conversation

---

### ### 8.2 AI Memory

- Stores:
  - Conversations
  - Exercises
  - Journal insights
  - Baseline profile
- Not user-editable
- Cannot be disabled
- Must never feel intrusive

---

## ## 9. CONVERSATION STATE MACHINE (AUTHORITATIVE)

### ### States

``

INIT  
CONVERSATIONAL\_DISCOVERY  
SUPPORTIVE\_PROCESSING  
EXERCISE\_SUGGESTION  
EXERCISE\_FACILITATION  
POST\_EXERCISE\_INTEGRATION  
CRISIS\_MODE  
``

### ### Rules

- Conversation always begins in `CONVERSATIONAL\_DISCOVERY`
- Exercises cannot be initiated without passing the **Exercise Suggestion Gate**
- CRISIS\_MODE overrides all other states

(See `conversationStateMachine.md` for full formal definition.)

---

## ## 10. THERAPEUTIC EXERCISE SYSTEM

### ### 10.1 Exercise Library

- 634 exercises
- 25 topics
- 13 therapeutic frameworks
- **No difficulty levels**

---

### ### 10.2 Exercise Data Model

```
``json
Exercise = {
 id: string,
 title: string,
 framework: string,
 topic: string,
 aspect: string,
 ai_prompt: string
}```
```

---

### ### 10.3 Exercise Discovery Flow (Conversation-First)

1. User engages in casual conversation
2. AI listens, reflects, validates
3. AI determines context sufficiency
4. AI suggests \*\*2–3 exercises max\*\*
5. User chooses:
  - Accept suggested exercise
  - Request custom exercise
  - Decline and continue conversation

---

### ### 10.4 Exercise Suggestion Gate (Mandatory)

AI may suggest exercises ONLY if:

1. A concrete challenge is articulated
2. The AI has reflected it accurately
3. The user is emotionally regulated
4. The AI can explain why structure helps
5. The user has not recently declined

---

### ### 10.5 Custom Exercises

- Must be based on existing frameworks
- Must follow framework phases exactly
- No new therapeutic methodologies allowed

---

### ### 10.6 Completion Criteria

- Reflection submission required
- AI confirms completion
- Progress logged

---

## ## 11. JOURNAL SYSTEM

### ### 11.1 Entry Types

- Free-form text
- Voice → transcription
- Media attachments (not analyzed)
- Prompt-based entries

---

### ### 11.2 AI Analysis

- Runs immediately on save (premium only)
- Modes:
  - Passive listening
  - Analytical processing
  - Therapeutic insights

---

### ### 11.3 Privacy

- Author-only access
- Never used for AI training
- Analysis visible only to user

---

## ## 12. ANALYTICS DASHBOARD

### ### Metrics

- Exercise acceptance vs decline
- Completion tracking
- Emotional trends
- Framework effectiveness
- Goal progress

### ### Access

- Separate analytics page
- Always available
- Optional engagement

---

## ## 13. GAMIFICATION SYSTEM

### ### Philosophy

Soft encouragement only. Non-competitive. Optional.

### ### Components

- 8 milestones
- 5 skill badges
- Points system
- Streaks (reset to zero on break)

---

## ## 14. UI SCREENS

- ### Core Screens
  - Authentication
  - Onboarding
  - AI Chat (Primary)
  - Exercise Flow (Inline)
  - Journal Editor
  - Analytics Dashboard
  - Gamification
  - Crisis Screen
  - Institutional Admin Dashboard

---

## ## 15. RESPONSIVE DESIGN

Breakpoint	Behavior
<640px	Single-column, chat-first
641–1024px	Expanded panels
>1024px	Multi-panel layout

---

## ## 16. TECHNICAL ARCHITECTURE

### ### Frontend

- Next.js (React)
- Tailwind CSS
- Zustand (global state)
- React Hook Form
- Recharts

### ### Backend

- Node.js
- REST APIs
- PostgreSQL
- Prisma ORM

### ### Auth

- NextAuth
- JWT sessions

### ### AI Layer

- Unified AI gateway
- Model routing by tier
- System prompt injection
- Memory persistence

### ### Payments

- Stripe subscriptions
- Stripe invoicing (institutions)

---

## ## 17. DATA EXPORT

Users may export:

- Full account data (JSON / CSV)

---

## ## 18. NON-GOALS (EXPLICITLY OUT OF SCOPE)

- Offline support
- Peer-to-peer chat
- Therapist dashboards
- Clinical diagnosis
- Medical treatment

---

## ## FINAL STATUS

This PRD is:

-  Internally consistent
-  Safety-aware
-  Monetization-ready
-  AI-agent executable
-  MVP-complete

---

## ### HANDOFF INSTRUCTION

Provide this PRD, along with:

- `sentientSelf\_systemPrompt\_v3.txt`
- `conversationStateMachine.md`

...to an AI coding agent.

The agent should implement **\*\*exactly as specified\*\***, without interpretation.

---

If you want next:

- A **\*\*build order checklist for the AI agent\*\***
- A **\*\*database schema\*\***
- Or a **\*\*one-command deployment plan\*\***

Just tell me.